

ISLR_CH8

Sky Liu


3/15/2019

8.4.1

```
par(xpd = NA)
plot(NA, NA, type = "n", xlim = c(0, 80), ylim = c(0, 100), xlab = "X", ylab = "Y")
# t1: x = 40; (40, 0) (40, 100)
lines(x = c(50, 50), y = c(0, 100))
text(x = 50, y = 108, labels = c("t1"), col = "red")
# t2: y = 75; (0, 75) (50, 75)
lines(x = c(0, 50), y = c(75, 75))
text(x = -8, y = 75, labels = c("t2"), col = "red")

# t3: x = 20; (20, 0) (20, 75)
lines(x = c(20, 20), y = c(0, 75))
text(x = 20, y = 80, labels = c("t3"), col = "red")

text(x = (40 + 75)/2, y = 50, labels = c("R1"))
text(x = 20, y = (100 + 75)/2, labels = c("R2"))
text(x = 30, y = 75/2, labels = c("R3"))
text(x = 10, y = 75/2, labels = c("R4"))
```



8.4.2

Boosting using depth-one trees is essentially equivalent to the additive model: for every tree generated, $\hat{f}^b(x)$ is fit with $d = 1$ split to the training data X , which amounts to single-variable linear model $\hat{f}_j(x)$. In total there are b_{max} trees, so there are b_{max} linear models generated.

Finally, the output of boosting is represented by


$$\hat{f}(x) = \sum_{b=1}^{b_{max}} \lambda \hat{f}^b(x)$$

, taking λ out of the sum, it can be written as a form of additive model.

$$\hat{f}(x) = \lambda \sum_{j=1}^{b_{max}} \hat{f}_j(x)$$

8.4.3

```
p <- seq(0, 1, 0.001)
gini <- 2*p*(1 - p)
entropy <- -(p * log(p) + (1 - p) * log(1 - p))
class.err <- 1 - pmax(p, 1 - p)
matplot(p, cbind(gini, entropy, class.err), col = c("red", "green", "blue"))
```



8.4.5

```
p <- c(0.1, 0.15, 0.2, 0.2, 0.55, 0.6, 0.6, 0.65, 0.7, 0.75)
```

majority vote approach

```
sum(p >= 0.5) < sum(p < 0.5)
```

```
## [1] FALSE
```

There are more red prediction, thus, it's RED

average approach

```
mean(p)
```

```
## [1] 0.45
```

The average predicted probability is $0.45 < 0.5$, thus, it's GREEN.


8.4.7

```
attach(Boston)
set.seed(1)
train <- sample(1:nrow(Boston), nrow(Boston)/2)
X.train <- Boston[train, -14]
X.test <- Boston[-train, -14]
Y.train <- Boston[train, 14]
Y.test <- Boston[-train, 14]

p <- dim(Boston)[2] - 1
p_2 <- p/2
p_sq <- sqrt(p)

rf_p <- randomForest(X.train, Y.train, xtest = X.test, ytest = Y.test,
  mtry = p, ntree = 500)
rf_p_2 <- randomForest(X.train, Y.train, xtest = X.test, ytest = Y.test,
  mtry = p_2, ntree = 500)
rf_p_sq <- randomForest(X.train, Y.train, xtest = X.test, ytest = Y.test,
  mtry = p_sq, ntree = 500)

plot(1:500, rf_p$test$mse, col = "green", type = "l", xlab = "Number of Trees",
  ylab = "Test MSE", ylim = c(10, 19));lines(1:500, rf_p_2$test$mse, col = "red", type = "l");lines(1
```



$m=p/2$ seems to provide the smallest test MSE. When the number of trees gets larger, this result become stable.

8.4.8

part a


```
attach(Carseats)
train <- sample(1:nrow(Carseats), nrow(Carseats)/2)
Carseats_train <- Carseats[train, ]
Carseats_test <- Carseats[-train, ]
```

part b

```
tree_carseats <- tree(Sales ~ ., data = Carseats_train)
summary(tree_carseats)
```

```
##
## Regression tree:
## tree(formula = Sales ~ ., data = Carseats_train)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "Age" "CompPrice" "Income"
## [6] "Advertising"
## Number of terminal nodes: 19
## Residual mean deviance: 2.306 = 417.4 / 181
## Distribution of residuals:
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## -4.26500 -0.90750 -0.04562 0.00000 1.01900 3.38000
```

```
plot(tree_carseats)
text(tree_carseats)
```



```

pred_carseats <- predict(tree_carseats, Carseats_test)
mse <- mean((Carseats_test$Sales - pred_carseats)^2)
mse

```

```
## [1] 4.704263
```



Test MSE is 4.71

part c

```

cv_carseats <- cv.tree(tree_carseats, FUN = prune.tree)
par(mfrow = c(1, 2))
plot(cv_carseats$size, cv_carseats$dev, type = "b")
plot(cv_carseats$k, cv_carseats$dev, type = "b")

```



The test MSE is the smallest when size is 9

```
pruned_carseats <- prune.tree(tree_carseats, best = 9)
pred_pruned <- predict(pruned_carseats, Carseats_test)
pruned_mse <- mean((Carseats_test$Sales - pred_pruned)^2)
pruned_mse
```

```
## [1] 4.52989
```

Pruning the tree does not improves the test MSE (from 4.71 to 5.45).

part d

```
bag_carseats <- randomForest(Sales ~ ., data = Carseats_train, mtry = 10, ntree = 500, importance = T)
bag_pred <- predict(bag_carseats, Carseats_test)
hag_mse <- mean((Carseats_test$Sales - bag_pred)^2)
hag_mse
```

```
## [1] 2.47614
```

```
importance(bag_carseats)
```

```
##           %IncMSE IncNodePurity
## CompPrice 20.255434    143.621120
## Income    8.030479     95.521905
## Advertising 9.113427     75.897281
## Population 1.011174     62.633198
## Price     55.679505    551.304273
## ShelfLoc  59.353315    507.781908
## Age       14.214033    156.726871
## Education  2.673099     55.170619
## Urban     -1.558460      8.917585
## US        0.271480      8.300157
```

Bagging improves the test MSE to 2.91. CompPrice,ShelveLoc and Price are the most important variables.

part e

```
rf_carseats <- randomForest(Sales ~ ., data = Carseats_train, mtry = 8, ntree = 500, importance = T)
rf_pred <- predict(rf_carseats, Carseats_test)
rf_mse<-mean((Carseats_test$Sales - rf_pred)^2)
rf_mse
```


```
## [1] 2.380419
```

```
importance(rf_carseats)
```

	%IncMSE	IncNodePurity
## CompPrice	18.08153539	146.805587
## Income	7.28442087	96.001096
## Advertising	8.43128558	75.156087
## Population	-0.76485895	67.419272
## Price	51.35108180	525.358133
## ShelveLoc	57.90822915	504.874241
## Age	13.94944627	175.527905
## Education	0.75967142	58.020287
## Urban	-1.40413862	9.504252
## US	-0.01818094	10.638819

Random forest improves the test MSE to 2.93. Advertising,ShelveLoc and Price are the most important variables.

```
p <- ncol(Carseats) - 1
errors <- c()
for ( d in seq(1, p) ) {
  rf_carseats <- randomForest(Sales~., data = Carseats, subset = train, mtry = d, ntree = 500)
  rf_pred <- predict(rf_carseats, Carseats_test)
  errors[d] <- mean((Carseats_test$Sales - rf_pred)^2)
}
plot(errors, col = "red", type = "b", lwd = 2, xlab = "m", ylab = "Test MSE")
```



```
which.min(errors)
```

```
## [1] 6
```

when $m = 8$, the test MSE is lowest

8.3.11

part a

```
attach(Caravan)
train = 1:1000
caravan <- Caravan
caravan$Purchase = ifelse(caravan$Purchase == "Yes", 1, 0)
caravan_train = caravan[train, ]
caravan_test = caravan[-train, ]
```

part b

```
set.seed(3)

boost_caravan <- gbm(Purchase~., data=caravan_train,
                     n.trees=1000, shrinkage=0.01, distribution="bernoulli", cv.folds = 10)

## Warning in gbm.fit(x = x, y = y, offset = offset, distribution =
## distribution, : variable 50: PVRAAUT has no variation.

## Warning in gbm.fit(x = x, y = y, offset = offset, distribution =
## distribution, : variable 71: AVRAAUT has no variation.
```



```
head(summary.gbm(boost_caravan, plotit = FALSE), 10)
```

```
##           var    rel.inf
## PPERSAUT PPERSAUT 14.822726
## MKOOPKLA MKOOPKLA  9.647725
## MOPLHOOG MOPLHOOG  6.516380
## MBERMIDD MBERMIDD  5.667377
## PBRAND    PBRAND   5.425667
## MGODGE    MGODGE   4.172987
## ABRAND    ABRAND   4.043386
## MINK3045  MINK3045  3.980106
## MOSTYPE   MOSTYPE  2.836807
## MSKA      MSKA     2.572698
```

part c

```
boost_prob = predict(boost_caravan, caravan_test, n.trees = 1000, type = "response")
boost_pred = ifelse(boost_prob > 0.2, 1, 0)
table(caravan_test$Purchase, boost_pred)
```

```
##    boost_pred
##      0      1
## 0 4413  120
## 1  257   32
```

```
32/(32+120)
```

```
## [1] 0.2105263
```

More than 20% of positive prediction are correct.

```
lm_caravan = glm(Purchase ~ ., data = caravan_train, family = binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
lm_prob = predict(lm_caravan, caravan_test, type = "response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```
lm_pred = ifelse(lm_prob > 0.2, 1, 0)
table(caravan_test$Purchase, lm_pred)
```

```
##    lm_pred
##      0      1
## 0 4183  350
## 1  231   58
```

```
58/(58+350)
```

```
## [1] 0.1421569
```

only 14% of positive prediction are correct. Boosting makes a better prediction.