

# ISLR\_CH9


*Sky Liu*

3/21/2019

## 9.7.3


part a

```
x1 = c(3, 2, 4, 1, 2, 4, 4)
x2 = c(4, 2, 4, 4, 1, 3, 1)
colors = c("red", "red", "red", "red", "blue", "blue", "blue")
plot(x1, x2, col = colors, xlim = c(0, 5), ylim = c(0, 5))
```



part b

```
plot(x1, x2, col = colors, xlim = c(0, 5), ylim = c(0, 5))
abline(-0.5, 1)
```




part c

$$-0.5 + X_1 - X_2 = 0.$$


part d

```
plot(x1, x2, col = colors, xlim = c(0, 5), ylim = c(0, 5))
abline(-0.5, 1)
abline(-1, 1, lty = 2)
abline(0, 1, lty = 2)
```



part e

```
plot(x1, x2, col = colors, xlim = c(0, 5), ylim = c(0, 5))
abline(-0.5, 1)
arrows(2, 1, 1.9, 1.4)
arrows(2, 2, 2.1, 1.6)
arrows(4, 4, 4.1, 3.6)
arrows(4, 3, 3.9, 3.4)
```




part f

because it's outside of the margin.

part g


```
plot(x1, x2, col = colors, xlim = c(0, 5), ylim = c(0, 5))
abline(-0.75, 1)
```



$$-0.75 + X_1 - X_2 = 0.$$

h

```
plot(x1, x2, col = colors, xlim = c(0, 5), ylim = c(0, 5))
points(c(3),c(1), col = c("red"))
```




### 9.7.5

part a

```
set.seed(3)
x1 = runif(500) - 0.5
x2 = runif(500) - 0.5
y = 1 * (x1^2 - x2^2 > 0)
```

part b

```
plot(x1[y == 0], x2[y == 0], col = "red", xlab = "X1", ylab = "X2", pch = "+")
points(x1[y == 1], x2[y == 1], col = "blue", pch = 4)
```




part c

```
lm_fit = glm(y ~ x1 + x2, family = binomial)
summary(lm_fit)
```

```
##
## Call:
## glm(formula = y ~ x1 + x2, family = binomial)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.257  -1.194   1.115   1.152   1.201
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.06632    0.08960   0.740   0.459
## x1          -0.06040    0.31237  -0.193   0.847
## x2          -0.20247    0.30630  -0.661   0.509
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 692.64  on 499  degrees of freedom
## Residual deviance: 692.16  on 497  degrees of freedom
## AIC: 698.16
##
## Number of Fisher Scoring iterations: 3
```

part d

```
data = data.frame(x1 = x1, x2 = x2, y = y)
lm_prob = predict(lm_fit, data, type = "response")
# probability threshold = 0.53
lm_pred = ifelse(lm_prob > 0.53, 1, 0)
data_pos = data[lm_pred == 1, ]
data_neg = data[lm_pred == 0, ]
plot(data_pos$x1, data_pos$x2, col = "blue", xlab = "X1", ylab = "X2", pch = "+")
points(data_neg$x1, data_neg$x2, col = "red", pch = 4)
```



part e


```
lm_fit = glm(y ~ poly(x1, 2) + poly(x2, 2) + I(x1 * x2), data = data, family = binomial)
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```


part f

```
lm_prob = predict(lm_fit, data, type = "response")
# probability threshold = 0.5
lm_pred = ifelse(lm_prob > 0.5, 1, 0)
data_pos = data[lm_pred == 1, ]
data_neg = data[lm_pred == 0, ]
plot(data_pos$x1, data_pos$x2, col = "blue", xlab = "X1", ylab = "X2", pch = "+")
points(data_neg$x1, data_neg$x2, col = "red", pch = 4)
```



part g


```
svm_fit = svm(as.factor(y) ~ x1 + x2, data, kernel = "linear", cost = 0.1)
svm_pred = predict(svm_fit, data)
data_pos = data[svm_pred == 1, ]
data_neg = data[svm_pred == 0, ]
plot(data_pos$x1, data_pos$x2, col = "blue", xlab = "X1", ylab = "X2", pch = "+")
points(data_neg$x1, data_neg$x2, col = "red", pch = 4)
```





part h

```
svm_fit = svm(as.factor(y) ~ x1 + x2, data, gamma = 1)
svm_pred = predict(svm_fit, data)
data_pos = data[svm_pred == 1, ]
data_neg = data[svm_pred == 0, ]
plot(data_pos$x1, data_pos$x2, col = "blue", xlab = "X1", ylab = "X2", pch = "+")
points(data_neg$x1, data_neg$x2, col = "red", pch = 4)
```



part i

Logistic regression with non-interactions and SVMs with linear kernels fail to find the decision boundary. Logistic regression with interactions find a decision boundary that is very close to the real one. But finding the correct interactions might become a challenge. However, using radial basis kernels, and only change the parameter gamma, saves lots of effort and simple CV could accomplish that. Thus, SVMs with non-linear kernel is very useful to find a non-linear decision boundary.

### 9.7.7

part a

```
gas.med = median(Auto$mpg)
new.var = ifelse(Auto$mpg > gas.med, 1, 0)
Auto$mpglevel = as.factor(new.var)
```

part b

linear kernel

```
set.seed(4)
tune_out = tune(svm, mpglevel ~ ., data = Auto, kernel = "linear", ranges = list(cost = c(0.01,
  0.1, 1, 5, 10, 100)))
summary(tune_out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.01275641
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-02 0.07378205 0.03437509
## 2 1e-01 0.05346154 0.03026685
## 3 1e+00 0.01275641 0.01344780
## 4 5e+00 0.02044872 0.01619554
## 5 1e+01 0.02038462 0.01074682
## 6 1e+02 0.03320513 0.01737168
```

When cost = 1, the test MSE is the smallest.

part c

polynomial kernal

```
set.seed(5)
tune_out = tune(svm, mpglevel ~ ., data = Auto, kernel = "polynomial", ranges = list(cost = c(0.1,
  1, 5, 10), degree = c(2, 3, 4)))
summary(tune_out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##    10      2
##
## - best performance: 0.5635897
##
## - Detailed performance results:
##   cost degree      error dispersion
## 1  0.1      2 0.5712179 0.04971286
## 2  1.0      2 0.5712179 0.04971286
## 3  5.0      2 0.5712179 0.04971286
## 4 10.0      2 0.5635897 0.05854295
## 5  0.1      3 0.5712179 0.04971286
```

```
## 6 1.0 3 0.5712179 0.04971286
## 7 5.0 3 0.5712179 0.04971286
## 8 10.0 3 0.5712179 0.04971286
## 9 0.1 4 0.5712179 0.04971286
## 10 1.0 4 0.5712179 0.04971286
## 11 5.0 4 0.5712179 0.04971286
## 12 10.0 4 0.5712179 0.04971286
```

When  $\text{cost} = 10$ ,  $\text{degree} = 2$ , the error is the smallest.

## radial kernel

```
set.seed(6)
tune_out = tune(svm, mpglevel ~ ., data = Auto, kernel = "radial", ranges = list(cost = c(0.1,
1, 5, 10), gamma = c(0.01, 0.1, 1, 5, 10, 100)))
summary(tune_out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   10 0.1
##
## - best performance: 0.02051282
##
## - Detailed performance results:
##   cost gamma      error dispersion
## 1 0.1 1e-02 0.09179487 0.04023204
## 2 1.0 1e-02 0.07141026 0.03778996
## 3 5.0 1e-02 0.04852564 0.03074258
## 4 10.0 1e-02 0.02301282 0.02549182
## 5 0.1 1e-01 0.07903846 0.03893119
## 6 1.0 1e-01 0.05371795 0.03716942
## 7 5.0 1e-01 0.02820513 0.03299190
## 8 10.0 1e-01 0.02051282 0.02911006
## 9 0.1 1e+00 0.52128205 0.14693828
## 10 1.0 1e+00 0.05878205 0.03646034
## 11 5.0 1e+00 0.05871795 0.03205442
## 12 10.0 1e+00 0.05871795 0.03205442
## 13 0.1 5e+00 0.55628205 0.05494470
## 14 1.0 5e+00 0.48987179 0.06983539
## 15 5.0 5e+00 0.48987179 0.06983539
## 16 10.0 5e+00 0.48987179 0.06983539
## 17 0.1 1e+01 0.55878205 0.05128210
## 18 1.0 1e+01 0.52557692 0.06532066
## 19 5.0 1e+01 0.50775641 0.06213538
## 20 10.0 1e+01 0.50775641 0.06213538
## 21 0.1 1e+02 0.55878205 0.05128210
## 22 1.0 1e+02 0.55878205 0.05128210
## 23 5.0 1e+02 0.55878205 0.05128210
```


```
## 24 10.0 1e+02 0.55878205 0.05128210
```

When cost = 10, gamma = 0.1, the error is the smallest.



part d

```
svm_linear = svm(mpglevel ~ ., data = Auto, kernel = "linear", cost = 1)
svm_poly = svm(mpglevel ~ ., data = Auto, kernel = "polynomial", cost = 10,
  degree = 2)
svm_radial = svm(mpglevel ~ ., data = Auto, kernel = "radial", cost = 10, gamma = 0.01)
plotpairs = function(fit) {
  for (name in names(Auto)[!(names(Auto) %in% c("mpg", "mpglevel", "name"))]) {
    plot(fit, Auto, as.formula(paste("mpg~", name, sep = "")))
  }
}
plotpairs(svm_linear)
```



**SVM classification plot**




**SVM classification plot**



horsepower  
**SVM classification plot**




**SVM classification plot**






```
plotpairs(svm_poly)
```




**SVM classification plot**




**SVM classification plot**



**SVM classification plot**




**SVM classification plot**





```
plotpairs(svm_radial)
```

**SVM classification plot**








**SVM classification plot**



**SVM classification plot**





## 9.7.8

part a

```
set.seed(7)
train = sample(dim(OJ)[1], 800)
OJ_train = OJ[train, ]
OJ_test = OJ[-train, ]
```

part b

```
svm_linear = svm(Purchase ~ ., kernel = "linear", data = OJ_train, cost = 0.01)
summary(svm_linear)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ_train, kernel = "linear",
##     cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##     cost: 0.01
##   gamma: 0.05555556
##
```

```
## Number of Support Vectors: 443
##
## ( 223 220 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM
```

Support vector classifier creates 443 support vectors from 800 training points. In the rest of points, 223 belong to level CH and 220 belong to level MM.

#### part c

```
train_pred = predict(svm_linear, OJ_train)
table(OJ_train$Purchase, train_pred)
```

```
##      train_pred
##           CH  MM
## CH 426  59
## MM  83 232
```

```
(59+83)/(426+232+59+83)
```

```
## [1] 0.1775
```

The training data error rate is 17.75%

```
test_pred = predict(svm_linear, OJ_test)
table(OJ_test$Purchase, test_pred)
```

```
##      test_pred
##           CH  MM
## CH 151  17
## MM  26  76
```

```
(17+26)/(151+17+26+76)
```

```
## [1] 0.1592593
```

The test data error rate is 16%

#### part d

```
set.seed(8)
tune_out = tune(svm, Purchase ~ ., data = OJ_train, kernel = "linear", ranges = list(cost = 10^seq(-2,1)
summary(tune_out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##      cost
```

```
## 3.162278
##
## - best performance: 0.17375
##
## - Detailed performance results:
##      cost    error dispersion
## 1  0.01000000 0.19250 0.03782269
## 2  0.01778279 0.18875 0.03508422
## 3  0.03162278 0.18750 0.03004626
## 4  0.05623413 0.18000 0.03238227
## 5  0.10000000 0.17750 0.03322900
## 6  0.17782794 0.17875 0.03120831
## 7  0.31622777 0.17750 0.03525699
## 8  0.56234133 0.17875 0.03634805
## 9  1.00000000 0.17500 0.03679900
## 10 1.77827941 0.17750 0.04073969
## 11 3.16227766 0.17375 0.04466309
## 12 5.62341325 0.17625 0.04101575
## 13 10.00000000 0.17625 0.03928617
```

The optimal cost is 3.1623

part e

```
svm_linear = svm(Purchase ~ ., kernel = "linear", data = OJ_train, cost = tune_out$best.parameters$cost)
train_pred = predict(svm_linear, OJ_train)
table(OJ_train$Purchase, train_pred)
```

```
##      train_pred
##      CH  MM
## CH 425  60
## MM  71 244
```

```
(60+71)/(60+71+425+244)
```

```
## [1] 0.16375
```

The training error rate after tuning is 16.375%

```
test_pred = predict(svm_linear, OJ_test)
table(OJ_test$Purchase, test_pred)
```

```
##      test_pred
##      CH  MM
## CH 152  16
## MM  26  76
```

```
(16+26)/(16+26+152+76)
```

```
## [1] 0.1555556
```

The test error rate after tuning is 15.56%

part f

```
set.seed(9)
svm_radial = svm(Purchase ~ ., data = OJ_train, kernel = "radial")
summary(svm_radial)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ_train, kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##         cost:  1
##        gamma: 0.05555556
##
## Number of Support Vectors: 368
##
## ( 183 185 )
##
##
## Number of Classes: 2
##
## Levels:
##  CH MM
```

Support vector classifier with radial kernel creates 368 support vectors from 800 training points. In the rest of points, 183 belong to level CH and 185 belong to level MM.

```
train_pred = predict(svm_radial, OJ_train)
table(OJ_train$Purchase, train_pred)
```

```
##      train_pred
##      CH  MM
## CH 438  47
## MM  82 233
```

```
(47+82)/(438+233+47+82)
```

```
## [1] 0.16125
```

The training error rate is 16.125%

```
test_pred = predict(svm_radial, OJ_test)
table(OJ_test$Purchase, test_pred)
```

```
##      test_pred
##      CH  MM
## CH 153  15
## MM  30  72
```

```
(15+30)/(153+72+15+30)
```

```
## [1] 0.1666667
```

The test error rate is 16.67%

```

set.seed(10)
tune_out = tune(svm, Purchase ~ ., data = OJ_train, kernel = "radial", ranges = list(cost = 10^seq(-2, 10, length.out = 100)))
#summary(tune_out)
svm_radial = svm(Purchase ~ ., data = OJ_train, kernel = "radial", cost = tune_out$best.parameters$cost)
train_pred = predict(svm_radial, OJ_train)
table(OJ_train$Purchase, train_pred)

```

```

##      train_pred
##      CH  MM
## CH 444  41
## MM  74 241

```

```

(41+74)/(41+74+444+241)

```

```

## [1] 0.14375

```

The training error rate after tuning is 14.375%

```

test_pred = predict(svm_radial, OJ_test)
table(OJ_test$Purchase, test_pred)

```

```

##      test_pred
##      CH  MM
## CH 152  16
## MM  30  72

```

```

(16+30)/(152+72+16+30)

```

```

## [1] 0.1703704

```

The test rate after tuning is 17%

## part g

```

set.seed(11)
svm_poly = svm(Purchase ~ ., data = OJ_train, kernel = "poly", degree = 2)
summary(svm_poly)

```

```

##
## Call:
## svm(formula = Purchase ~ ., data = OJ_train, kernel = "poly",
##      degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:  1
##   degree:  2
##   gamma:  0.05555556
##   coef.0:  0
##
## Number of Support Vectors:  441
##
## ( 223 218 )
##

```

```
##
## Number of Classes: 2
##
## Levels:
## CH MM
```

Support vector classifier with radial kernel creates 441 support vectors from 800 training points. In the rest of points, 183 belong to level CH and 223 belong to level 218

```
train_pred = predict(svm_poly, OJ_train)
table(OJ_train$Purchase, train_pred)
```

```
##      train_pred
##      CH  MM
## CH 449  36
## MM 111 204
```

```
(36+111)/(36+111+449+204)
```

```
## [1] 0.18375
```

The training error rate is 18.375%

```
test_pred = predict(svm_poly, OJ_test)
table(OJ_test$Purchase, test_pred)
```

```
##      test_pred
##      CH  MM
## CH 158  10
## MM  41  61
```

```
(10+41)/(10+41+158+61)
```

```
## [1] 0.1888889
```

The test error rate is 18.89%

```
set.seed(12)
tune_out = tune(svm, Purchase ~ ., data = OJ_train, kernel = "poly", degree = 2, ranges = list(cost = 10^0:10^4))
#summary(tune_out)
svm_poly = svm(Purchase ~ ., data = OJ_train, kernel = "poly", cost = tune_out$best.parameters$cost)
train_pred = predict(svm_poly, OJ_train)
table(OJ_train$Purchase, train_pred)
```

```
##      train_pred
##      CH  MM
## CH 446  39
## MM  80 235
```

```
(39+80)/(39+80+446+235)
```

```
## [1] 0.14875
```

The training error rate after tuning is 14.875%

```
test_pred = predict(svm_poly, OJ_test)
table(OJ_test$Purchase, test_pred)
```

```
##      test_pred
##      CH  MM
## CH 156  12
```

```
## MM 32 70
```

```
(12+32)/(12+32+156+70)
```

```
## [1] 0.162963
```

The test rate after tuning is 16.3%

### part h

With linear kernel, the training error rate after tuning is 16.375% and the test error rate after tuning is 15.56%

With radial kernel, the training error rate after tuning is 14.375% and the test error rate after tuning is 17%

With polynomial kernel, the training error rate after tuning is 14.875% and the test error rate after tuning is 16.3%

Suprisingly, on this data, linear kernel has smaller test error rate.